

# Scikit learn

Wednesday, 13 November 2024 6:26 PM

- **fit**: Trains the model or transformer on your data.
- **predict**: Makes predictions on new data, used with estimators like classifiers or regressors.
- **transform**: Applies a transformation.
- **Pipeline**: Combines multiple steps, making code cleaner and ensuring consistent data processing across all stages.

## Using fit and predict with a Classifier

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.datasets import load_iris

# Load data
X, y = load_iris(return_X_y=True)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Fit the model
clf = LogisticRegression()
clf.fit(X_train, y_train)

# Predict
predictions = clf.predict(X_test)
print(predictions)
```

## Explanation

- `train_test_split`: Used to split the dataset into training and testing sets.
- `LogisticRegression`: The classifier used for our model.
- `load_iris`: Loads the Iris dataset, a commonly used dataset for classification.
  - `X, y = load_iris(return_X_y=True)`
- Here, `X` is the feature data (input), and `y` is the target data (output).
- `return_X_y=True` gives us `X` and `y` directly.
  - `X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)`
- The data is split into a training set (70%) and a test set (30%).
- `random_state=42` ensures reproducibility.

- `clf = LogisticRegression()`

We instantiate a `LogisticRegression` model and store it in `clf`.

- `clf.fit(X_train, y_train)`
- `fit` trains the model using `X_train` (features) and `y_train` (target).
- During training, the model learns to classify based on patterns in the data.

- `predictions = cir.predict(X_test)`
- We use `predict` on the `X_test` data to generate predictions.
- `predictions` is an array with predicted class labels for each sample in `X_test`.
  - `print(predictions)`
- Displays the predictions, which are the model's classification of the `X_test` samples.

#### Using fit and transform with a Transformer

```
from sklearn.preprocessing import StandardScaler
import numpy as np

# Data
data = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])

# Fit and transform
scaler = StandardScaler()
scaler.fit(data) # Learns the mean and std
transformed_data = scaler.transform(data) # Transforms the data
print(transformed_data)
```

#### Import Necessary Libraries:

- `StandardScaler` is used for standardizing features by removing the mean and scaling to unit variance.

`data = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])`

- We create a 2D array, `data`, to demonstrate scaling.

`scaler = StandardScaler()`

`StandardScaler` is initialized, which will later calculate the mean and standard deviation.

`scaler.fit(data)`

`fit` calculates the mean and standard deviation for each feature column.

`transformed_data = scaler.transform(data)`

- `transform` scales the data using the previously calculated mean and standard deviation.
- `transformed_data` now contains the standardized version of data.

`print(transformed_data)`

Displays the transformed data.

#### Using a Pipeline to Chain Steps Together

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split

# Load data
X, y = load_iris(return_X_y=True)
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Create a pipeline
pipeline = Pipeline([
    ('scaler', StandardScaler()),      # Step 1: Scale features
    ('classifier', LogisticRegression()) # Step 2: Fit logistic regression
])

# Fit pipeline on training data
pipeline.fit(X_train, y_train)

# Predict with pipeline
predictions = pipeline.predict(X_test)
print(predictions)

```

- Pipeline from `sklearn.pipeline` lets us chain together transformations and a model.
- We also import `StandardScaler`, `LogisticRegression`, and the dataset utilities as in previous examples.

We load the Iris dataset and split it into training and testing sets as before.

```

pipeline = Pipeline([
    ('scaler', StandardScaler()),      # Step 1: Scale features
    ('classifier', LogisticRegression()) # Step 2: Fit logistic regression
])

```

- We create a `Pipeline` object that contains two steps:
- `scaler`: Uses `StandardScaler` to standardize the data.
- `classifier`: Uses `LogisticRegression` as the classifier.
- Each step has a name (`'scaler'`, `'classifier'`) and an instance of a transformer or model.

```
pipeline.fit(X_train, y_train)
```

- `fit` applies each step in sequence:
- First, it scales `X_train` using `StandardScaler`.
- Then, it trains the `LogisticRegression` model on the scaled data.
- The pipeline manages the transformations and model fitting automatically, ensuring consistent processing.

```
predictions = pipeline.predict(X_test)
```

When we call `predict`, the pipeline applies the same scaling transformation to `X_test` and then makes predictions with the `LogisticRegression` model.

```
print(predictions)
```

Displays the predictions, just as before, but this time done through the pipeline.

Key Benefits of Using Pipelines

- **Consistency:** Ensures the same preprocessing steps are applied to both training and test

data.

- Simplified Code: Reduces code complexity by chaining multiple steps.
- Hyperparameter Tuning: Pipelines allow easy integration with GridSearchCV and RandomizedSearchCV for hyperparameter tuning across multiple steps.