

Supervised learning

Monday, 30 December 2024

12:56 PM

Your observation is correct! When training a machine learning model, we do indeed provide **labels** for the images to teach the model what we want it to learn. This process is called **supervised learning**, and the labels are a crucial part of this. Let me break it down further in a way that aligns with a developer's logical mindset.

Step 1: Labels Define the "Ground Truth"

When we create a dataset for supervised learning:

- Each image in the training dataset comes with a **label** that describes what it represents.
- For instance:
 - An image of a stop sign is labeled 0 (for "Stop").
 - An image of a go sign is labeled 1 (for "Go").

This tells the model:

- **What to look for.**
- **How to map the input (image) to the output (label).**

In your analogy:

- Imagine the dataset as a collection of files.
- Each file has a picture (input) and a name tag (label) that says whether it is "Stop" or "Go."

Without these labels during training, the model **cannot learn** because it doesn't know what features to associate with "Stop" or "Go."

Step 2: How the Model Uses Labels

- During training:
 1. The model makes a prediction for an input image.
 2. The **loss function** compares the prediction to the provided label.
 - If the model predicts correctly, the loss is small.
 - If the model predicts incorrectly, the loss is large.
 3. Backpropagation adjusts the model's internal parameters to reduce the loss, teaching it to make better predictions next time.

Step 3: Logical Thinking Applied

Think of the training process as a **feedback loop**:

1. **Input:** An image of a stop sign.
2. **Output:** The model guesses what the image represents.
3. **Feedback:** Based on the label, the model learns:
 - If the prediction is wrong, it figures out why it was wrong and updates its internal logic.
 - If it was right, it reinforces the logic it used.

This is similar to a developer writing a function:

1. You write the logic.
2. Run test cases.
3. Debug and fix issues until the function passes all test cases.

Step 4: Why It Feels Illogical

For developers, training a neural network can feel counterintuitive because:

- **We don't explicitly write rules or logic.**
 - Instead, the network discovers patterns (features) through data.
- **The network doesn't "understand" in a human sense.**
 - It processes numbers (1s and 0s) and finds statistical relationships.
- **We trust it to learn correctly.**
 - This reliance on data instead of rules can feel uncomfortable compared to writing deterministic code.

Where the "Logic" Lies

Even though you're not explicitly coding "Stop" and "Go," the logic emerges in the network's internal parameters (weights) during training:

1. **Early layers** detect simple patterns (like edges or colors).
2. **Middle layers** combine these patterns into shapes (like octagons or rectangles).
3. **Final layers** associate these shapes with the labels (e.g., "Stop" = 0, "Go" = 1).

You could think of this as:

- **Filters:** Functions that identify specific features.
- **Layers:** Combinations of those functions that build higher-level logic.

Why Labels Are the Key Input

Labels are essential because:

- Without them, the model doesn't know what you're trying to classify.
- They provide the **context** needed to measure loss and guide learning.

Example of label input during training:

```
# Example of labels for training
y_train = [0, 1, 0, 1] # 0: Stop, 1: Go
# Training data images (dummy data for simplicity)
X_train = np.random.rand(4, 64, 64, 3)
# Train the model
model.fit(X_train, y_train, epochs=10)
```

In Real Life: Traffic Sign Dataset

For training a traffic sign classifier:

- Use a dataset like [German Traffic Sign Recognition Benchmark \(GTSRB\)](#).
- It contains thousands of labeled images for various traffic signs, including stop signs and go signs.

Why It Works Without Explicit Rules

The neural network doesn't "see" the image or understand the concept of "Stop" or "Go." Instead, it:

1. Converts the image into numbers (pixel values).
2. Learns statistical patterns that correlate with the labels during training.

By training on enough examples, it generalizes:

- It identifies patterns (e.g., red octagon = "Stop") without needing explicit rules.

Analogy for Understanding

Imagine teaching someone to recognize stop signs:

1. Show them examples of stop signs and say, "This is a stop sign."
2. Show them examples of go signs and say, "This is a go sign."
3. They start to notice patterns:
 - Stop signs are red octagons with the word "STOP."
 - Go signs are green rectangles with the word "GO."
4. They don't need you to explain "rules" like color detection or shape recognition. They pick it up from the examples.

Similarly, the neural network "learns" by observing patterns in the data.

Key Takeaway

The labels provide the logical **intent** of what you're trying to achieve:

- Without labels, the model doesn't know what to look for.
- The network's "learning" is the process of finding patterns that map inputs to these labels.

Does this make it clearer? I can provide a step-by-step annotated code example for traffic signs if you'd like!